

Getting good performance from your application

Tuning techniques for serial programs on cache-based computer systems

Overview

- ❑ Introduction
- ❑ Memory Hierarchy
- ❑ General Optimization Techniques
- ❑
- ❑ Compilers
- ❑ Analysis Tools
- ❑ Tuning Guide

De-vectorization

- ❑ Cache space and bandwidth are scarce resources
- ❑ Compilers know this – but sometimes they have to store data that does not need to be stored.
- ❑ This impacts:
 - ❑ bandwidth
 - ❑ cache capacity
 - ❑ instruction scheduling

De-vectorization

- ❑ A typical problem with scratch data stored in vectors
- ❑ Difficult/impossible for the compiler to detect
- ❑ Depends on coding style

De-vectorization – Example

```
COMMON /SCRATCH/TMP (N)
```

```
DO I = 1, N
  TMP(I) = ...
  :
  ... = TMP(I)
END DO
```

```
...
```

```
DO I = 1, N
  TMP(I) = ..
  :
END DO
```

- *Because TMP() is global, the compiler has to store it in the first loop*
- *In the second loop, TMP() is overwritten, but the compiler will most likely not see this*
- *The programmer may know that TMP() is a scratch array only*

De-vectorization – Solutions

Array TMP needed later on:

```
REAL T1, TMP (N)
```

```
DO I = 1, N
  T1 = ...
  :
  ... = T1
END DO
```

```
...
```

```
DO I = 1, N
  TMP(I) = ..
  :
END DO
```

Array TMP not needed later on:

```
REAL T1
```

```
DO I = 1, N
  T1 = ...
  :
  ... = T1
END DO
```

```
...
```

```
DO I = 1, N
  T1 = ..
  :
END DO
```

Stripmining

- Large loops are difficult to optimize
- Especially the *register allocation in the compiler has a hard time and can get confused*
- Splitting the loop into smaller loops may improve performance
- However, this may cause scalars (local to the loop) to be replaced by vectors
- On very large loops this will increase memory usage notably
- Through *stripmining* memory usage can be kept under control

```
DO I = 1, LONG
  X(I) = ...
  A = ...
  Y(I) = A + ...
END DO
```

Split loop in two parts

```
DO I = 1, LONG
  X(I) = ...
  VA(I) = ...
END DO
```

```
DO I = 1, LONG
  Y(I) = VA(I) + ...
END DO
```

Stripmining – Code structure

```
DO i = is, ie
  .....
END DO
```

Stripmine this loop with length NS

```
DO i1 = is, ie, ns
  ivl = min(ie-i1+1, ns)
  DO i2 = 0, ivl-1
    i = i1+i2
    .....
  END DO
END DO
```

Split inner loop

```
DO i1 = is, ie, ns
  ivl = min(ie-i1+1, ns)
  DO i2 = 0, ivl-1
    i = i1+i2
    .....
  END DO
  DO i2 = 0, ivl-1
    i = i1+i2
    .....
  END DO
  DO i2 = 0, ivl-1
    i = i1+i2
    .....
  END DO
END DO
```

Best practice

It is up to you to write code such that the compiler can find opportunities for optimization:

- ❑ Write efficient, but clear code
- ❑ Avoid very “fat” (bulky) loops
- ❑ Design your data structures carefully
- ❑ Minimize global data

Best practice

- ❑ Branches:
 - ❑ simplify where possible
 - ❑ try to split the branch part out of the loop
- ❑ Avoid function calls in loops (use inlining)
- ❑ Leave the low level details to the compiler

Summary

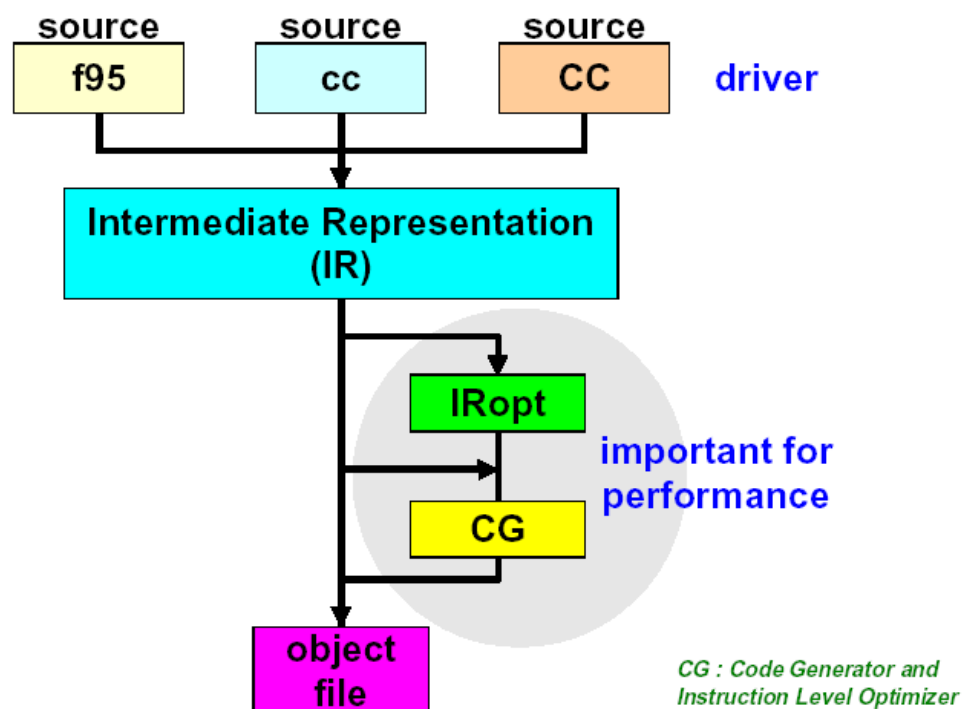
- ❑ Most tuning techniques presented here are generic, i.e. they (probably/hopefully) improve your code on all cache based systems.
- ❑ The *tuning parameters* may be different, though, since they depend on the underlying hardware:
 - ❑ cache sizes and levels
 - ❑ prefetchand your problem's *memory footprint*
- ❑ Use the *best* compiler available on your platform.

The Oracle Solaris Studio Compilers
(a.k.a. Sun Studio)
—
an example how compilers work

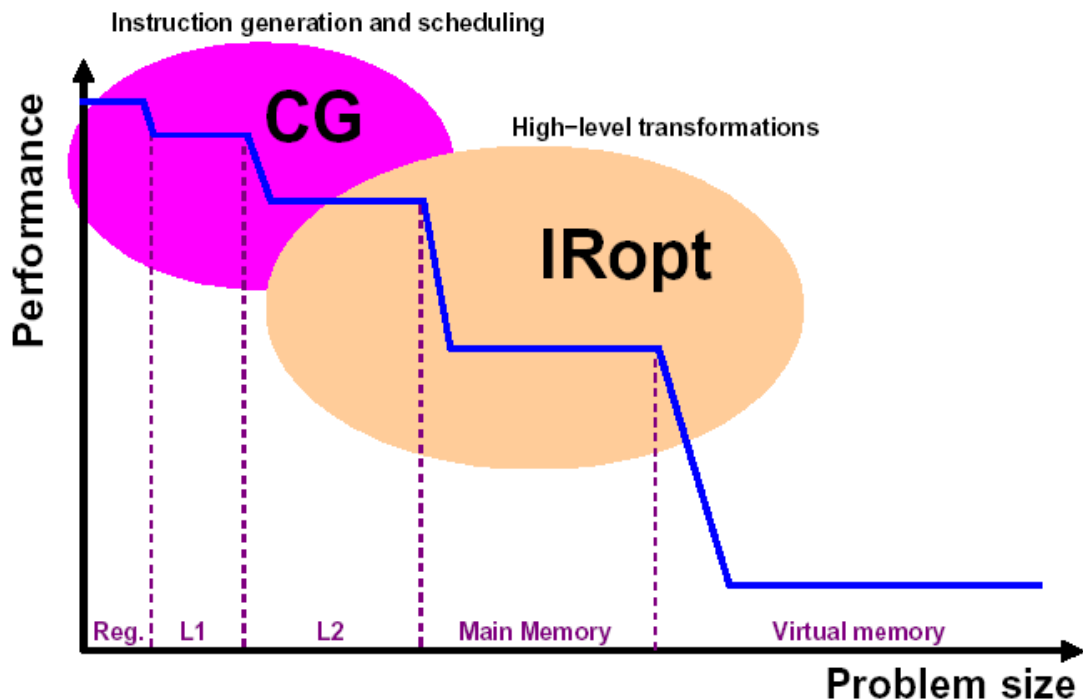
Sun Studio: Overview

- ❑ Compiler Components
 - ❑ Compiler Options
 - ❑ Compiler Commentary
 - ❑ Best practice
- ❑ Note: to use the Studio tools, you have to load the module first:
'module load studio'

Sun Studio: Compiler Components



Sun Studio: Who does what?



Sun Studio: Minimal Compiler Options

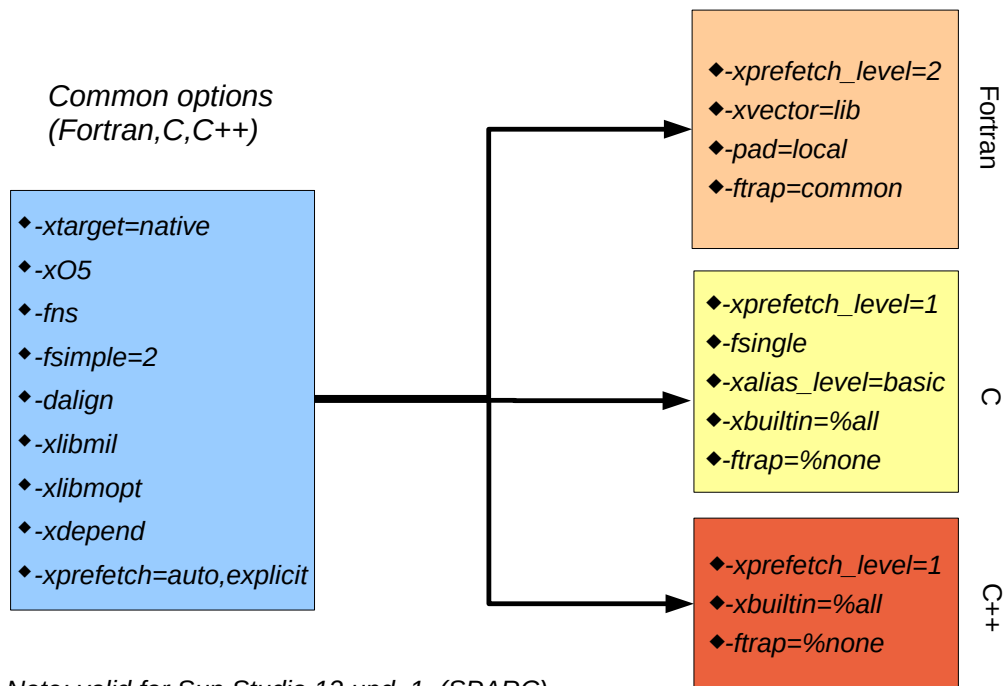
In general, one gets very good performance by just using 2 options for compiling and linking:

`-g -fast`

For specific x86_64-processors (cross-compiling):

<i>AMD Opteron:</i>	<code>-g -fast -xchip=opteron</code>
<i>Intel Nehalem:</i>	<code>-g -fast -xchip=nehalem</code>
<i>Intel Westmere:</i>	<code>-g -fast -xchip=westmere</code>
<i>Intel Sandy Bridge:</i>	<code>-g -fast -xchip=sandybridge</code>
<i>Intel Ivy Bridge:</i>	<code>-g -fast -xchip=ivybridge</code>
<i>Intel Haswell:</i>	<code>-g -fast -xchip=haswell</code>
<i>Intel Broadwell:</i>	<code>-g -fast -xchip=broadwell</code>
<i>Intel Skylake:</i>	<code>-g -fast -xchip=skylake</code>

Sun Studio: The -fast macro



Note: valid for Sun Studio 12 upd. 1 (SPARC)

Sun Studio: Recommendations

- ❑ Use a Makefile and `make` for compiling/linking:

```
OPT      = -g -fast
ISA      =
CHIP     =
CFLAGS   = $(OPT) $(ISA) $(CHIP)
```

- ❑ Always start with *-fast* !
- ❑ The compilers follow the 'rightmost option wins' rule, i.e. one can overrule options defined by the -fast macro.

Sun Studio: Recommendations

- ❑ `-fast` is a convenience macro that (in general) gives optimal performance with one single option
- ❑ `-fast` can change from one release to another!
- ❑ Use '`-fast -xdryrun`' to check what `-fast` expands to

What about other compilers? - I

- ❑ Other compilers have similar options, that combine many optimizations into a single option
- ❑ GCC: `-O`, `-O3`
- ❑ Intel: `-O2` (default!), `-O3`, `-fast`
- ❑ look up in the manpages/documentation, what that corresponds to
- ❑ my experience: difficult to find out, not so easy to switch off 'unwanted options'

What about other compilers? - II

- ❑ GCC: show optimizer options
 - ❑ `gcc -Q -help=optimizers -O3`
- ❑ shows a list of all known '-f...' options and their status
- ❑ switch extra options on/off, e.g. loop unrolling
 - ❑ `-funroll-loops`
 - ❑ `-fno-unroll-loops`

What about other compilers? - III


- ❑ Other compilers work in a similar way, but ...
- ❑ I've chosen Oracle Studio for various reasons:
 - ❑ my own experience
 - ❑ it is easy to use
 - ❑ it makes it easy to test effects of certain options
 - ❑ it is a nice tool for teaching
 - ❑ it makes it easier for you to understand certain aspects in this class
- ❑ Downside: It seems that Oracle has stopped further development of the product in 2017!

Sun Studio: Optimization levels

- ❑ The level of optimization can be specified with the `-xOn` option
- ❑ Each level includes the lower levels
- ❑ Choices for `n`:
 - ❑ `n = 1` – Basic block level optimizations
 - ❑ `n = 2` – Some additional optimizations
 - ❑ `n = 3` – Loop transformations and modulo scheduling
 - ❑ `n = 4` – Intra-file inlining and pointer tracking
 - ❑ `n = 5` – Aggressive optimizations


Inter Procedural Optimization

- ❑ With the `-xipo` option, the compiler stores additional information into the object files
- ❑ This information is used during the link phase to perform additional optimizations



Prefetch: Hiding memory latency

- ❑ The number of clock cycles to access memory increases with the CPU clock speed.
- ❑ Prefetch is a way to overcome this:
 - Fetch data ahead in time, anticipating future use.
- ❑ Special prefetch instructions must be available



Prefetch Support


- ❑ Prefetch is a common feature in modern CPUs: both data and instruction prefetch.
- ❑ Implementation is system dependent!
- ❑ Examples of CPUs that support prefetch:
 - ❑ AMD Opteron
 - ❑ Intel (since Pentium 3, 4)
 - ❑ UltraSPARC IIIi, IIICu, IV, IV+
- ❑ Compilers have to support prefetch.
- ❑ There is hardware prefetch, too (Xeon, Opteron)

Sun Studio: Prefetch options

- ❑ Automatic – done by the compiler:
 - ❑ -xprefetch=yes
 - ❑ Control level and type:
 - ❑ -xprefetch_level=n (n = 1, 2 or 3)
 - ❑ -xprefetch_auto_type=[no%]indirect_array_access
- ❑ Explicit – controlled by the user:
 - ❑ Either through functions calls (C/C++) or directives (Fortran).
 - ❑ Can be combined with automatic prefetch.
 - ❑ For more information see the compiler manual.

Prefetch: example

Example: Matrix times vector in C (row version)



US-IIIi @ 1062 MHz
 L1 : 64 kB
 L2 : 1 MB
 Peak : 2.1 Gflop/s

Pointer overlap – or “aliasing”

```
void vecadd(int n, double *a, double *b, double *c)
{
    for(int i = 0; i < n; i++)
        c[i] = a[i] + b[i];
}
```

```
vecadd(n, &a[0], &b[0], &a[1]);
```

```
void vecadd(n, &a[0], &b[0], &a[1])
{
    for(int i = 0; i < n; i++)
        a[i+1] = a[i] + b[i];
}
```

Data
dependency!

Sun Studio: C code and -xrestrict


- ❑ Pointer aliasing problem: The C compiler has to assume that different pointers may overlap:
 - ❑ Correct – but non-optimal – code will be generated
 - ❑ Only the programmer might know, that there is no overlap.
- ❑ You can tell the compiler that there is no overlap, with the `-xrestrict` option.
- ❑ Note: It is then your responsibility that this assumption will not be violated!

Sun Studio: C code and -xrestrict

- ❑ Syntax: `-xrestrict=n`, where `n` is one of
 - ❑ `%none` : all pointers may overlap (default)
 - ❑ `%all` : no pointers overlap
 - ❑ `f1[,f2]` : a comma separated list of function names
- ❑ `-xrestrict` is the same as `-xrestrict=%all`
- ❑ Warning: Use `-xrestrict` only if you can make sure that the pointers never overlap!

Sun Studio: C code and -xrestrict

Vector addition example: non-overlapping vectors!



US-III @ 900 MHz
 L1 : 64 kB
 L2 : 8 MB
 Peak : 1.8 Gflop/s
 Sun Studio 10

Pointer overlap – fix your code

- ❑ If you can assure that the pointers don't overlap, you can fix your code using the C99 'restrict' keyword:

```
void  
vecadd(int n, double * restrict a,  
       double * restrict b, double * restrict c)  
{  
    for(int i = 0; i < n; i++)  
        c[i] = a[i] + b[i];  
}
```

- ❑ Advantage: local change, no global assumption
- ❑ Needs a C99 compliant compiler to be portable!

Sun Studio: Useful options – 1

- ❑ -flags : Lists all the available compiler flags on the screen (long list)
- ❑ -xhelp=readme : Displays the README file (release notes) on the screen
- ❑ -xdryrun : see what the compiler would do (macro expansion, no compilation!)
- ❑ -V : Shows the compiler version

```
% f90 -V  
f90: Sun Fortran 95 8.4 SunOS_sparc 2009/06/03  
Usage: f90 [ options ] files. Use 'f90 -flags' for details
```

Sun Studio: Useful options – 2

- ❑ -g : Generates debugging information and adds compiler commentary to the object file. Necessary if the performance analysis tools should be used.
- ❑ **Note:** -g can be used in connection with optimization!
- ❑ -xlibmopt : Optimized version of libm
- ❑ -xlibmil: Use fast assembly language inline templates for standard functions
- ❑ -xvector : vectorized intrinsics, e.g. sqrt()

Sun Studio: Useful options – 3

- ❑ -xprefetch_level=n : Try different levels of prefetching ($n = 1, 2, 3$)
- ❑ -xrestrict : (C/C++) tell the compiler that your pointers do point to 'restricted areas' in memory, i.e. no overlap of arrays
- ❑ -stackvar : (Fortran) keep local data on stack
- ❑ -xlic_lib=sunperf : Link with the Sun Performance Library (LAPACK, BLAS, Sparse-BLAS (Netlib and NIST), FFT, Sparse Solvers)

Sun Studio: Useful options – 4

- ❑ Options that are useful under development – but should be avoided in production:
 - ❑ -C : array boundary checking (Fortran)
 - ❑ -pg/-p : Unix profilers – you don't really need them, use the Performance Analyzer instead.
 - ❑ -xcheck=... : implements runtime checks.
 - ❑ -xcheck=init_local : initialize local variables with a value that is likely to cause an arithmetic exception (Fortran)
 - ❑ -xcheck=stkovf : check for stack overflow.
 - ❑ -Xlist : check your Fortran code for inconsistencies (commented source in .lst file)

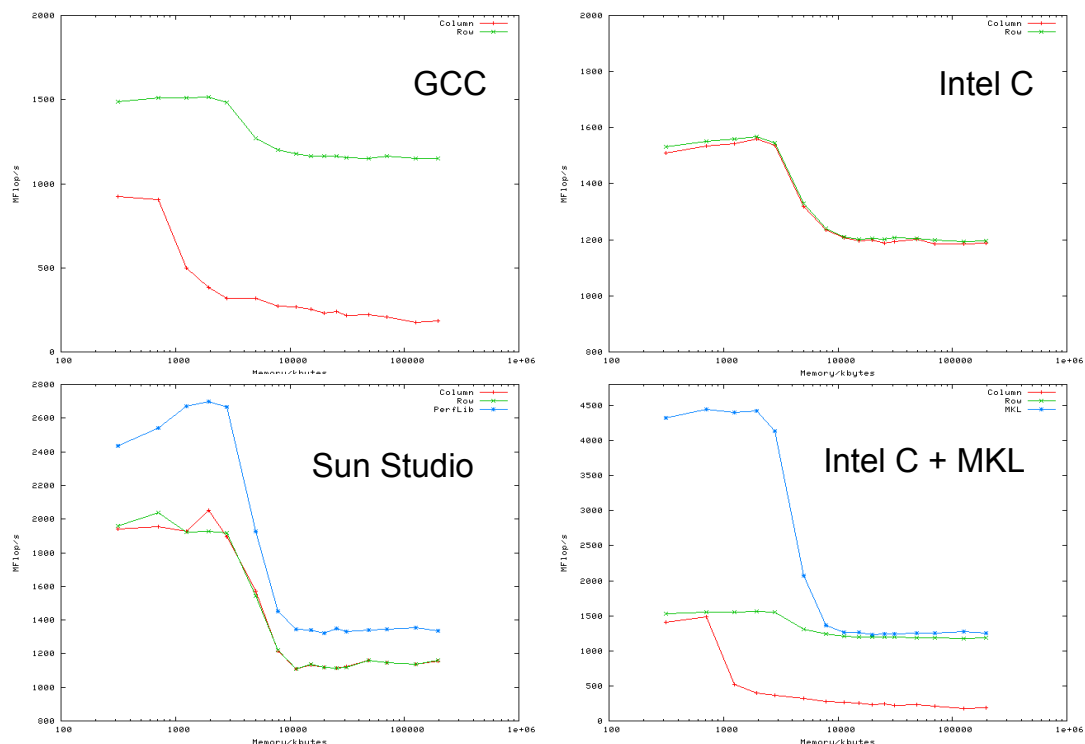
Sun Studio: Where to get it?

- ❑ Go to [Oracle TechNetwork](#)
- ❑ Download and install it – it is for free!
- ❑ Best practice: It is always good to have more than one compiler (or compiler version) at hand!
 - ❑ to test compatibility / portability
 - ❑ to exclude compiler bugs

Matrix times vector


- ❑ 2 code versions: **Column** and **Row**
- ❑ 3 compilers under Linux:
 - ❑ GCC
 - ❑ Intel C
 - ❑ Sun Studio C
- ❑ 2 libraries:
 - ❑ Intel MKL
 - ❑ Sun Performance Library
- ❑ Note: this is a quick comparison! (June 2008)

Matrix times vector



Matrix times vector

Comparison for large data sets:



January 2019

02614 - High-Performance Computing

117

Ease of use – vendor library

❑ Intel/GCC with Intel MKL:

- ❑ several compiler options
- ❑ several linker options
- ❑ depends on platform: IA32 or EMT64
- ❑ set up of run-time environment
- ❑ different downloads, installations and licenses

❑ Sun Studio:

- ❑ one linker option: `-xlic_lib=sunperf`
- ❑ no need to set up run-time environment
- ❑ one download, installation – and no license

January 2019

02614 - High-Performance Computing

118

Compiler commentary

Sun Studio: Compiler Commentary

How do I know what the compiler did with my code?

- ❑ Compile with -g (C++: -g0) and use the `er_src` command on the object files:
 - ❑ This generates a code listing with comments from the compiler (also used in the Performance Analyzer GUI).
- ❑ Command: `er_src -src foo.c foo.o`
- ❑ Note: the examples on the next slides are from the Solaris SPARC version of Sun Studio – the x64 versions can differ

Sun Studio: Compiler Commentary

```
// ex.c
#define A(i,j) a[(i)*n + (j)]
#define B(i,j) b[(i)*n + (j)]
#define C(i,j) c[(i)*n + (j)]

void
mysub(int m, int n, double *a, double *b,
      double *c) {

    int i, j;

    for(j = 0; j < n; j++)
        for(i = 0; i < m; i++)
            C(i,j) = A(i,j) + B(i,j);
}
```

Sun Studio: Compiler Commentary

```
5. void
6. mysub(int m, int n, double *a,
        double *b, double *c){
    <Function: mysub>
7.
8.     int i, j;
9.
```

```
% cc -g -fast -c ex.c
% er_src -src ex.c ex.o
```

UltraSPARC

Source loop below has tag L1

```
10.     for(j = 0; j < n; j++)
```

Source loop below has tag L2

L2 scheduled with steady-state cycle count = 5

L2 unrolled 2 times

L2 has 2 loads, 1 stores, 6 prefetches, 1 FPadds, 0
FPMuls, and 0 FPdivs per iteration

```
11.         for(i = 0; i < m; i++)
12.             C(i,j) = A(i,j) + B(i,j);
13.     }
```

Sun Studio: Compiler Commentary

```

5. void
6. mysub(int m, int n, double *a,
        double *b, double *c){
    <Function: mysub>
7.
8.     int i, j;
9.
Source loop below has tag L1
L1 interchanged with L2
L1 scheduled with steady-state cycle count = 2
L1 unrolled 4 times
L1 has 2 loads, 1 stores, 3 prefetches, 1 FPaddd, 0
FPmul, and 0 FPdivs per iteration

10.    for(j = 0; j < n; j++)
Source loop below has tag L2
L2 interchanged with L1
11.        for(i = 0; i < m; i++)
12.            C(i,j) = A(i,j) + B(i,j);
13. }

```

January 2019

02614 - High-Performance Computing

123

UltraSPARC

Sun Studio: Compiler Commentary

```

5. void
6. mysub(int m, int n, double *a,
        double *b, double *c){
    <Function: mysub>
7.
8.     int i, j;
9.
Source loop below has tag L1
L1 interchanged with L2
L1 cloned for microvectorizing-epilog. Clone is L6
L1 is micro-vectorized

10.    for(j = 0; j < n; j++)
Source loop below has tag L2
L2 interchanged with L1
11.        for(i = 0; i < m; i++)
12.            C(i,j) = A(i,j) + B(i,j);
13. }

```

January 2019

02614 - High-Performance Computing

124

x86_64

Sun Studio: Compiler Commentary

```
% cc -g -fast -xrestrict -xvector=no%simd -c ex.c
% er_src -src ex.c ex.o
```

```
5. void
6. mysub(int m, int n, double *a,
        double *b, double *c){
    <Function: mysub>
7.
8.     int i, j;
9.
```

x86_64

Source loop below has tag L1
L1 interchanged with L2

```
10.     for(j = 0; j < n; j++)
```

Source loop below has tag L2
L2 interchanged with L1

```
11.         for(i = 0; i < m; i++)
12.             C(i,j) = A(i,j) + B(i,j);
13. }
```

Sun Studio: Compiler Commentary

```
% cc -g -fast -xrestrict -c ex2.c
% er_src -src ex2.c ex2.o
```

```
7. void
8. mysub(int m, int n, double *a,
        double *b, double *c){
    <Function: mysub>
9.
10.     int i, j;
11.
12.     for(j = 0; j < n; j++)
13.         for(i = 0; i < m; i++) {
14.             C(i,j) = A(i,j) + B(i,j);
```

Function mean not inlined because the compiler has
not seen the body of the routine

```
15.         mean(C(i,j));
16.     }
17. }
```

Compiler Commentary – I

Why is this useful?

- ❑ Compilers are no longer a black box!
- ❑ What the compiler has done – or hasn't done/couldn't do – to the code is made visible to the programmer.
- ❑ Useful information is provided, so a programmer can take action, e.g.
 - ❑ code changes
 - ❑ different set of compiler options

Compiler Commentary – II

More reasons, why it is really useful:

- ❑ annotations are where they belong: in the code – and not on the screen during compilation (or in a log file)
- ❑ review at any time – even a long time after the code has been compiled
- ❑ visible in the analyzer output as well, together with runtime profile

Other tricks

Reconstruct the compiler options from the object files and/or executable:

- ❑ `dwarfdump file.o` (C/C++/Fortran)
- ❑ `dumpstabs file.o` (Fortran, before Studio 12u1)
- ❑ and look for
 - ❑ `command_line` (`dwarfdump`) or
 - ❑ `CMDLINE` (`dumpstabs`) in the output.
- ❑ Very useful to check what has been done.
- ❑ Note: `dwarfdump` works for GCC `.o` files, too

Analysis tools

Analysis tools

- ❑ analysis tools are useful to detect bottlenecks in codes
- ❑ modern analysis tools (unlike “old” profilers) work even on 'non-instrumented' code: no need to recompile (in principle)
- ❑ runtime profiles down to the source level (profilers usually work on function/subroutine level)

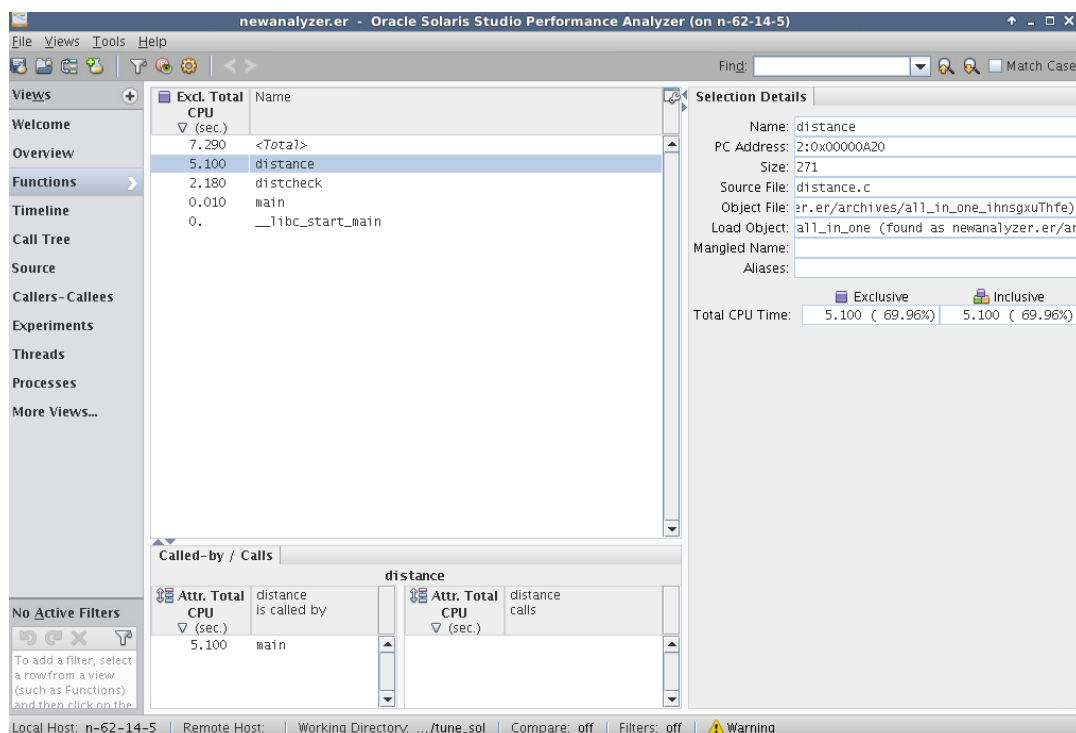
Analysis tools

- ❑ Oracle: Solaris Studio Performance Analyzer
 - ❑ Linux (x64)
- ❑ Intel: Vtune Performance Analyzer (Windows/Linux)
- ❑ Mac OS X: Instruments (part of Xcode)
- ❑ 'perf' command line tool (Linux)
- ❑ Google Performance Tools (Linux/Windows):
 - ❑ collection of runtime libraries and command line tools

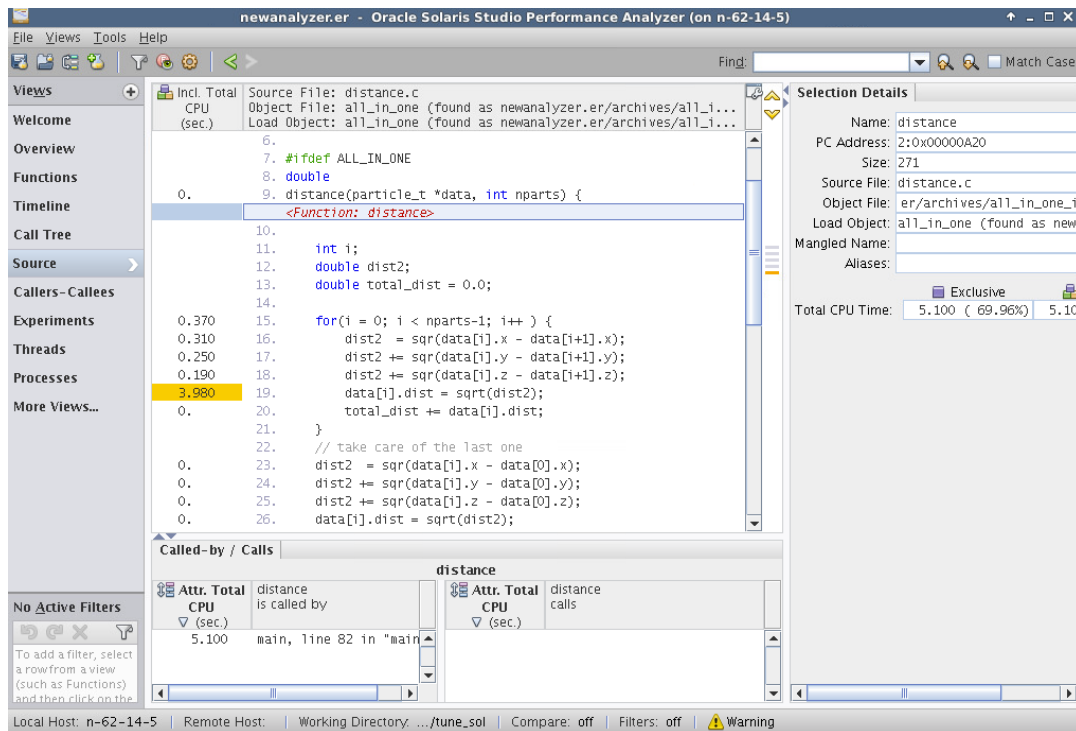
Sun Studio: Performance Analyzer

- ❑ Sun Studio provides a powerful toolset for runtime analysis
- ❑ Both GUI and command line tools
 - ❑ analyzer – GUI for collecting and analyzing performance data
 - ❑ collect – Command to collect performance data
 - ❑ er_print – Command to analyze performance data in ASCII format (good for scripting)

Sun Studio: Performance Analyzer



Sun Studio: Performance Analyzer



Hardware Performance Counters

- ❑ Almost all modern CPUs have built-in hardware performance counters:
 - ❑ How many instructions were executed?
 - ❑ How many clock cycles were used?
 - ❑ How many L1 data cache misses occurred?
- ❑ The supported counters are usually listed in the architecture reference manuals.
- ❑ Be aware: The counter names are not for beginners!

Using the Performance Counters

- ❑ Native OS tools, e.g. Linux:
 - ❑ perf – Performance monitoring tool
 - ❑ requires newer Linux kernel (> 2.6.31)
 - ❑ examples:


```
% perf stat -e <event_name> -- command
% perf stat -e <event_name> -p PID

% perf record -e <event_name> -- command
% perf report
```
 - ❑ 'perf top' - requires root privileges

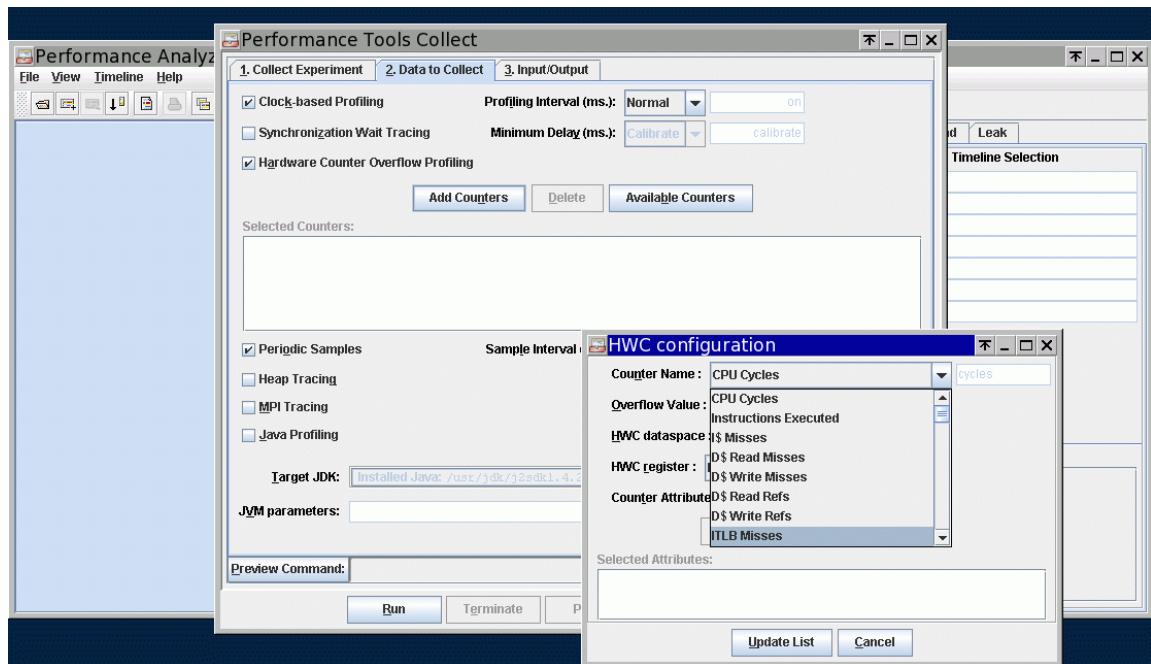
Using the Performance Counters

- ❑ Available performance counters:
 - ❑ system and CPU dependent
 - ❑ get a list:
 - ❑ % perf list
 - ❑ % collect (no argument)
 - ❑ example: no. of available performance counters on

❑ US-IV:	70
❑ US-IV+:	101
❑ AMD Opteron:	169
❑ Xeon E5-... v3:	269

Using the Performance Counters

Activating performance counters in analyzer:



January 2019

02614 - High-Performance Computing

141

Sun Studio: Performance Analyzer

Analyzer demo

January 2019

02614 - High-Performance Computing

142

Tuning Guide – compact version

- ❑ Make a 'baseline' version (with different data sets/memory requirements)
- ❑ Try to find the best compiler options
 - ❑ with or w/o prefetching
 - ❑ ...
- ❑ Use analysis tools to locate the 'hot spots'
- ❑ Introduce code changes
- ❑ Repeat the last two steps until you are satisfied

End of lecture 2